

```
In [1]: # importing required libraries
%matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import style
import seaborn as sns
```

## Week 1

```
In [2]: # Load data set for processing and analysis
data = pd.read_csv("C:\MySpace\Data_Science\Capstone\Project2\health care diabetes.csv")
```

```
In [3]: # Data mining to understand the data and take necessary action w.r.t. data correction
```

```
In [4]: data.head()
```

```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [5]: data.describe()
```

```
Out[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

In [6]:

```
data.corr()
```

Out[6]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>Pregnancies</b>	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341	0.221898
<b>Glucose</b>	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514	0.466581
<b>BloodPressure</b>	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528	0.065068
<b>SkinThickness</b>	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970	0.074752
<b>Insulin</b>	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163	0.130548
<b>BMI</b>	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242	0.292695
<b>DiabetesPedigreeFunction</b>	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.033561	0.173844
<b>Age</b>	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1.000000	0.238356
<b>Outcome</b>	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.173844	0.238356	1.000000

In [7]:

```
#required Column to Update for Null (0) Values
Sel_Col=['Glucose','BloodPressure','SkinThickness','Insulin','BMI']
Sel_Col
```

Out[7]: ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']

```
In [8]: # Count of zero in fields ,Interested columns are ['Glucose','BloodPressure','SkinThickness','Insulin','BMI']
Col_Zero_Val=data[data[Sel_Col]==0].count()
Col_Zero_Val
```

```
Out[8]: Pregnancies      0
Glucose      5
BloodPressure 35
SkinThickness 227
Insulin      374
BMI          11
DiabetesPedigreeFunction 0
Age          0
Outcome      0
dtype: int64
```

```
In [9]: #Updating required column with their respective Median value
# We are not considering Mean as there are rows present with "0" values and due to same mean values are not the correct value

for i in data.columns:
    if (i in Sel_Col):
        m=data[data[i]!=0][i].median()
#         print(i)
#         print(m)
        data[i]=data[i].apply(lambda x: m if x==0 else x)

data.head()
```

```
Out[9]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148.0	72.0	35.0	125.0	33.6	0.627	50	1
1	1	85.0	66.0	29.0	125.0	26.6	0.351	31	0
2	8	183.0	64.0	29.0	125.0	23.3	0.672	32	1
3	1	89.0	66.0	23.0	94.0	28.1	0.167	21	0
4	0	137.0	40.0	35.0	168.0	43.1	2.288	33	1

```
In [10]: # ReChecking post updating null value with median
Col_Zero_Val=data[data[Sel_Col]==0].count()
```

Col\_Zero\_Val

```
Out[10]: Pregnancies      0
          Glucose         0
          BloodPressure   0
          SkinThickness   0
          Insulin         0
          BMI             0
          DiabetesPedigreeFunction  0
          Age             0
          Outcome         0
          dtype: int64
```

```
In [11]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Pregnancies           768 non-null   int64
 1   Glucose               768 non-null   float64
 2   BloodPressure         768 non-null   float64
 3   SkinThickness         768 non-null   float64
 4   Insulin               768 non-null   float64
 5   BMI                  768 non-null   float64
 6   DiabetesPedigreeFunction 768 non-null   float64
 7   Age                  768 non-null   int64
 8   Outcome              768 non-null   int64
dtypes: float64(6), int64(3)
memory usage: 54.1 KB
```

```
In [12]: Positive = data[data['Outcome']==1]
          Positive.head(5)
```

```
Out[12]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148.0	72.0	35.0	125.0	33.6	0.627	50	1
2	8	183.0	64.0	29.0	125.0	23.3	0.672	32	1

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>4</b>	0	137.0	40.0	35.0	168.0	43.1	2.288	33	1
<b>6</b>	3	78.0	50.0	32.0	88.0	31.0	0.248	26	1
<b>8</b>	2	197.0	70.0	45.0	543.0	30.5	0.158	53	1

In [13]: `data[Sel_Col].value_counts()`

Out[13]:

Glucose	BloodPressure	SkinThickness	Insulin	BMI	
44.0	62.0	29.0	125.0	25.0	1
56.0	56.0	28.0	45.0	24.2	1
129.0	86.0	20.0	270.0	35.1	1
	90.0	7.0	326.0	19.6	1
	92.0	49.0	155.0	36.4	1
					..
105.0	100.0	36.0	125.0	43.3	1
106.0	52.0	29.0	125.0	31.2	1
	54.0	21.0	158.0	30.9	1
	56.0	27.0	165.0	29.0	1
199.0	76.0	43.0	125.0	42.9	1

Length: 768, dtype: int64

In [14]: `data['Glucose'].value_counts().head(7)`

Out[14]:


99.0	17
100.0	17
117.0	16
129.0	14
125.0	14
106.0	14
111.0	14

Name: Glucose, dtype: int64

In [15]: `plt.hist(data['Glucose'])`

Out[15]:

```
(array([ 4., 19., 87., 149., 166., 125., 88., 54., 44., 32.]),
 array([ 44. , 59.5, 75. , 90.5, 106. , 121.5, 137. , 152.5, 168. ,
        183.5, 199. ]),
 <BarContainer object of 10 artists>)
```




```
In [16]: data['BloodPressure'].value_counts().head(7)
```

```
Out[16]: 72.0    79
70.0    57
74.0    52
78.0    45
68.0    45
64.0    43
80.0    40
Name: BloodPressure, dtype: int64
```

```
In [17]: plt.hist(data['BloodPressure'])
```

```
Out[17]: (array([ 3.,  2., 35., 118., 261., 214., 105., 18., 10.,  2.]),
 array([ 24. , 33.8, 43.6, 53.4, 63.2, 73. , 82.8, 92.6, 102.4,
        112.2, 122. ]),
 <BarContainer object of 10 artists>)
```




```
In [18]: data['SkinThickness'].value_counts().head(7)
```

```
Out[18]: 29.0    244
32.0     31
30.0     27
27.0     23
23.0     22
28.0     20
33.0     20
Name: SkinThickness, dtype: int64
```

```
In [19]: plt.hist(data['SkinThickness'])
```

```
Out[19]: (array([ 59., 141., 408., 118., 36., 4., 1., 0., 0., 1.]),
array([ 7. , 16.2, 25.4, 34.6, 43.8, 53. , 62.2, 71.4, 80.6, 89.8, 99. ]),
<BarContainer object of 10 artists>)
```




```
In [20]: data['Insulin'].value_counts().head(7)
```

```
Out[20]: 125.0    378
105.0     11
130.0      9
140.0      9
120.0      8
 94.0       7
180.0       7
Name: Insulin, dtype: int64
```

```
In [21]: plt.hist(data['Insulin'])
```

```
Out[21]: (array([142., 517., 55., 29., 7., 10., 4., 1., 2., 1.]),
 array([ 14. , 97.2, 180.4, 263.6, 346.8, 430. , 513.2, 596.4, 679.6,
        762.8, 846. ]),
 <BarContainer object of 10 artists>)
```






```
In [22]: data['BMI'].value_counts().head(7)
```

```
Out[22]: 32.3    14
32.0    13
31.6    12
31.2    12
33.3    10
32.4    10
30.1     9
Name: BMI, dtype: int64
```

```
In [23]: plt.hist(data['BMI'])
```

```
Out[23]: (array([ 52., 161., 207., 193.,  91.,  48.,  10.,   4.,   1.,   1.]),
 array([18.2 , 23.09, 27.98, 32.87, 37.76, 42.65, 47.54, 52.43, 57.32,
        62.21, 67.1 ]),
 <BarContainer object of 10 artists>)
```



In [24]: `data.describe().transpose()`


Out[24]:

	count	mean	std	min	25%	50%	75%	max
<b>Pregnancies</b>	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.00000	17.00
<b>Glucose</b>	768.0	121.656250	30.438286	44.000	99.75000	117.0000	140.25000	199.00
<b>BloodPressure</b>	768.0	72.386719	12.096642	24.000	64.00000	72.0000	80.00000	122.00
<b>SkinThickness</b>	768.0	29.108073	8.791221	7.000	25.00000	29.0000	32.00000	99.00
<b>Insulin</b>	768.0	140.671875	86.383060	14.000	121.50000	125.0000	127.25000	846.00
<b>BMI</b>	768.0	32.455208	6.875177	18.200	27.50000	32.3000	36.60000	67.10
<b>DiabetesPedigreeFunction</b>	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.62625	2.42
<b>Age</b>	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.00000	81.00
<b>Outcome</b>	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.00000	1.00

## Week 2

In [25]: `plt.hist(Positive['BMI'],histtype='stepfilled',bins=20)`

```
Out[25]: (array([ 8., 10., 23., 41., 45., 40., 29., 19., 14., 17., 9., 4., 3.,
        3., 1., 0., 1., 0., 0., 1.]),
         array([22.9 , 25.11, 27.32, 29.53, 31.74, 33.95, 36.16, 38.37, 40.58,
        42.79, 45.   , 47.21, 49.42, 51.63, 53.84, 56.05, 58.26, 60.47,
        62.68, 64.89, 67.1 ]),
         [<matplotlib.patches.Polygon at 0x1cb83ca0>])
```




```
In [26]: Positive['BMI'].value_counts().head(7)
```

```
Out[26]: 32.9    8
         31.6    7
         33.3    6
         31.2    5
         30.5    5
         32.3    5
         32.0    5
         Name: BMI, dtype: int64
```

```
In [27]: plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20)
```

```
Out[27]: (array([ 3., 4., 5., 11., 14., 16., 20., 21., 20., 17., 15., 17., 14.,
        15., 15., 14., 12., 12., 10., 13.]),
         array([ 78.   , 84.05, 90.1 , 96.15, 102.2 , 108.25, 114.3 , 120.35,
        126.4 , 132.45, 138.5 , 144.55, 150.6 , 156.65, 162.7 , 168.75,
        174.8 , 180.85, 186.9 , 192.95, 199.   ]),
         [<matplotlib.patches.Polygon at 0x1cbb5bb0>])
```




```
In [28]: Positive['Glucose'].value_counts().head(7)
```

```
Out[28]: 125.0    7
128.0    6
158.0    6
129.0    6
115.0    6
146.0    5
181.0    5
Name: Glucose, dtype: int64
```

```
In [29]: plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20)
```

```
Out[29]: (array([ 1.,  0.,  1.,  0.,  6.,  5.,  3., 17., 25., 35., 68., 30., 25.,
        23., 14.,  4.,  3.,  3.,  2.,  3.]),
 array([ 30.,  34.2,  38.4,  42.6,  46.8,  51.,  55.2,  59.4,  63.6,
        67.8,  72.,  76.2,  80.4,  84.6,  88.8,  93.,  97.2, 101.4,
        105.6, 109.8, 114. ]),
 [<matplotlib.patches.Polygon at 0x1cbf6cd0>])
```




```
In [30]: Positive['BloodPressure'].value_counts().head(7)
```

```
Out[30]: 72.0    32
70.0    23
76.0    18
78.0    17
74.0    17
82.0    13
64.0    13
Name: BloodPressure, dtype: int64
```

```
In [31]: plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20)
```

```
Out[31]: (array([ 1.,  5., 11., 21., 113., 41., 34., 20., 15.,  4.,  1.,
        0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  1.]),
 array([ 7. , 11.6, 16.2, 20.8, 25.4, 30. , 34.6, 39.2, 43.8, 48.4, 53. ,
        57.6, 62.2, 66.8, 71.4, 76. , 80.6, 85.2, 89.8, 94.4, 99. ]),
 [<matplotlib.patches.Polygon at 0x1cc314a8>])
```




```
In [32]: Positive['SkinThickness'].value_counts().head(7)
```

```
Out[32]: 29.0    95
          32.0    14
          30.0     9
          33.0     9
          36.0     8
          37.0     8
          39.0     8
          Name: SkinThickness, dtype: int64
```

```
In [33]: plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20)
```

```
Out[33]: (array([ 4., 12., 165., 31., 18., 10., 8., 5., 2., 1., 1.,
           6., 2., 1., 1., 0., 0., 0., 0., 1.]),
          array([ 14. , 55.6, 97.2, 138.8, 180.4, 222. , 263.6, 305.2, 346.8,
           388.4, 430. , 471.6, 513.2, 554.8, 596.4, 638. , 679.6, 721.2,
           762.8, 804.4, 846. ]),
          [<matplotlib.patches.Polygon at 0x1cc65ac0>])
```



```
In [34]: Positive['Insulin'].value_counts().head(7)
```


```
Out[34]: 125.0    140
130.0      6
180.0      4
175.0      3
156.0      3
185.0      2
225.0      2
Name: Insulin, dtype: int64
```

```
In [35]: #Scatter plot
```


```
In [36]: BloodPressure = Positive['BloodPressure']
Glucose = Positive['Glucose']
SkinThickness = Positive['SkinThickness']
Insulin = Positive['Insulin']
BMI = Positive['BMI']
```

```
In [37]: plt.scatter(BloodPressure, Glucose, color=['b'])
plt.xlabel('BloodPressure')
plt.ylabel('Glucose')
```

```
plt.title('BloodPressure & Glucose')  
plt.show()
```




```
In [38]: g = sns.scatterplot(x= "Glucose" ,y= "BloodPressure",  
                           hue="Outcome",  
                           data=data);
```






```
In [39]: B =sns.scatterplot(x= "BMI" ,y= "Insulin",  
                        hue="Outcome",  
                        data=data);
```



```
In [40]: S =sns.scatterplot(x= "SkinThickness" ,y= "Insulin",  
                        hue="Outcome",  
                        data=data);
```




```
In [41]: ### correlation matrix
data.corr()
```

```
Out[41]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>Pregnancies</b>	1.000000	0.128213	0.208615	0.081770	0.025047	0.021559	-0.033523	0.544341	0.221898
<b>Glucose</b>	0.128213	1.000000	0.218937	0.192615	0.419451	0.231049	0.137327	0.266909	0.492782
<b>BloodPressure</b>	0.208615	0.218937	1.000000	0.191892	0.045363	0.281257	-0.002378	0.324915	0.165723
<b>SkinThickness</b>	0.081770	0.192615	0.191892	1.000000	0.155610	0.543205	0.102188	0.126107	0.214873
<b>Insulin</b>	0.025047	0.419451	0.045363	0.155610	1.000000	0.180241	0.126503	0.097101	0.203790
<b>BMI</b>	0.021559	0.231049	0.281257	0.543205	0.180241	1.000000	0.153438	0.025597	0.312038
<b>DiabetesPedigreeFunction</b>	-0.033523	0.137327	-0.002378	0.102188	0.126503	0.153438	1.000000	0.033561	0.173844
<b>Age</b>	0.544341	0.266909	0.324915	0.126107	0.097101	0.025597	0.033561	1.000000	0.238356
<b>Outcome</b>	0.221898	0.492782	0.165723	0.214873	0.203790	0.312038	0.173844	0.238356	1.000000


```
In [42]: ### create correlation heat map
sns.heatmap(data.corr())
```

```
Out[42]: <AxesSubplot:>
```




```
In [43]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True,cmap='viridis') ### gives correlation value
```

```
Out[43]: <AxesSubplot:>
```



```
In [44]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True) ### gives correlation value
```

```
Out[44]: <AxesSubplot:>
```



## Week 3

```
In [45]: # Logistic Regrestion and model building
```

In [46]: `data.head(5)`

Out[46]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148.0	72.0	35.0	125.0	33.6	0.627	50	1
1	1	85.0	66.0	29.0	125.0	26.6	0.351	31	0
2	8	183.0	64.0	29.0	125.0	23.3	0.672	32	1
3	1	89.0	66.0	23.0	94.0	28.1	0.167	21	0
4	0	137.0	40.0	35.0	168.0	43.1	2.288	33	1

In [47]: `features = data.iloc[:,[0,1,2,3,4,5,6,7]].values`  
`label = data.iloc[:,8].values`

In [48]: `pip install -U scikit-learn`

Requirement already satisfied: scikit-learn in c:\users\prabh\appdata\local\programs\python\python38-32\lib\site-packages (1.0.2)  
 Requirement already satisfied: numpy>=1.14.6 in c:\users\prabh\appdata\local\programs\python\python38-32\lib\site-packages (from scikit-learn) (1.21.4)  
 Requirement already satisfied: joblib>=0.11 in c:\users\prabh\appdata\local\programs\python\python38-32\lib\site-packages (from scikit-learn) (1.1.0)  
 Requirement already satisfied: scipy>=1.1.0 in c:\users\prabh\appdata\local\programs\python\python38-32\lib\site-packages (from scikit-learn) (1.7.3)  
 Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\prabh\appdata\local\programs\python\python38-32\lib\site-packages (from scikit-learn) (3.1.0)  
 Note: you may need to restart the kernel to use updated packages.  
 WARNING: You are using pip version 21.3.1; however, version 22.0.3 is available.  
 You should consider upgrading via the 'c:\users\prabh\appdata\local\programs\python\python38-32\python.exe -m pip install --upgrade pip' command.

In [49]: `#Train test split`  
`from sklearn.model_selection import train_test_split`  
`X_train,X_test,y_train,y_test = train_test_split(features,`  
 `label,`  
 `test_size=0.2,`  
 `random_state =10)`

```
In [50]: #Create model
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train,y_train)
```

c:\users\prabh\appdata\local\programs\python\python38-32\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
 n\_iter\_i = \_check\_optimize\_result(  
 LogisticRegression())

```
Out[50]:
print(model.score(X_train,y_train))
print(model.score(X_test,y_test))
```

```
0.7833876221498371
0.7272727272727273
```

```
In [52]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(label,model.predict(features))
cm
```

```
Out[52]: array([[445,  55],
               [120, 148]], dtype=int64)
```

```
In [53]: from sklearn.metrics import classification_report
print(classification_report(label,model.predict(features)))
```

	precision	recall	f1-score	support
0	0.79	0.89	0.84	500
1	0.73	0.55	0.63	268
accuracy			0.77	768
macro avg	0.76	0.72	0.73	768

weighted avg      0.77      0.77      0.76      768


## Week 4

```
In [54]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

AUC: 0.838

Out[54]: [





```
In [55]: #Applying Decission Tree Classifier  
from sklearn.tree import DecisionTreeClassifier  
model3 = DecisionTreeClassifier(max_depth=5)  
model3.fit(X_train,y_train)
```

```
Out[55]: DecisionTreeClassifier(max_depth=5)
```

```
In [56]: model3.score(X_train,y_train)
```

```
Out[56]: 0.8192182410423453
```

```
In [57]: model3.score(X_test,y_test)
```

```
Out[57]: 0.7467532467532467
```

```
In [58]: #Applying Random Forest  
from sklearn.ensemble import RandomForestClassifier  
model4 = RandomForestClassifier(n_estimators=11)  
model4.fit(X_train,y_train)
```

```
Out[58]: RandomForestClassifier(n_estimators=11)
```

```
In [59]: model4.score(X_train,y_train)
```

```
Out[59]: 0.99185667752443
```

```
In [60]: model4.score(X_test,y_test)
```

```
Out[60]: 0.7272727272727273
```

```
In [61]: #Support Vector Classifier  
  
from sklearn.svm import SVC  
model5 = SVC(kernel='rbf',
```

```
gamma='auto')
model5.fit(X_train,y_train)
```

Out[61]: SVC(gamma='auto')

```
In [62]: # model.score(X_test,y_test).score(X_train,y_train)
```

```
In [63]: model5.score(X_test,y_test)
```

Out[63]: 0.6168831168831169

```
In [64]: #Applying K-NN
from sklearn.neighbors import KNeighborsClassifier
model2 = KNeighborsClassifier(n_neighbors=7,
                             metric='minkowski',
                             p = 2)
model2.fit(X_train,y_train)
```

Out[64]: KNeighborsClassifier(n\_neighbors=7)

```
In [65]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score


# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".format(tpr,fpr,thresholds))
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

```
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
```

AUC: 0.848

True Positive Rate - [0. 0.07835821 0.25746269 0.4738806 0.64552239 0.7761194  
0.9141791 0.98880597 1. ], False Positive Rate - [0. 0.002 0.012 0.05 0.144 0.274 0.45 0.644 1. ] Thresholds - [2.  
1. 0.85714286 0.71428571 0.57142857 0.42857143  
0.28571429 0.14285714 0. ]  
Text(0, 0.5, 'True Positive Rate')

Out[65]:



In [66]:


```
#Precision Recall Curve for Logistic Regression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
```

```
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.'))
```

f1=0.628 auc=0.715 ap=0.716

Out[66]: [



In [67]:

```
#Precision Recall Curve for KNN


from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model2.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
```

```

# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.674 auc=0.771 ap=0.730  
 Out[67]: [<matplotlib.lines.Line2D at 0x1f5254c0>]



In [68]: *#Precision Recall Curve for Decision Tree Classifier*

```

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model3.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model3.predict(features)

```


```

# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.685 auc=0.791 ap=0.762

Out[68]: [



In [69]:

```

#Precision Recall Curve for Random Forest

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model4.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]

```

```
# predict class values
yhat = model4.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.909 auc=0.963 ap=0.955

Out[69]: [

